

Intro To dplyr

Steve Pittard

2020-01-24

Contents

1	R Data Structures	5
1.1	The Mighty Data Frame	5
1.2	A Reference Data Frame	6
1.3	Relation to dplyr	7
2	Digging In	9
2.1	Structure	11
2.2	Meta Information	11
2.3	Printing	12
2.4	Accessing Rows And Columns	12
2.5	Interrogation	14
2.6	Missing values	16
2.7	Continuous vs Factors	19
2.8	Sorting data	20
2.9	Creating Data Frames From CSV Files	22
3	Literature	23
4	Methods	25

Chapter 1

R Data Structures

There are a number of data structures in R such as **vectors**, **lists**, **matrices** and **arrays** but the premier data structure in R is known as the **data.frame**. This structure can be described as follows:

- A data frame is a special type of list that contains data in a format that allows for easier manipulation, reshaping, and open-ended analysis
- Data frames are tightly coupled collections of variables. It is one of the more important constructs you will encounter when using R so learn all you can about it
- A data frame is an analogue to the Excel spreadsheet but is much more flexible for storing, manipulating, and analyzing data
- Data frames can be constructed from existing vectors, lists, or matrices. Many times they are created by reading in comma delimited files, (CSV files), using the `read.table` command

Once you become accustomed to working with data frames, R becomes so much easier to use. In fact, it could be well argued tht UNTIL you wrap your head around the data frame concept then you cannot be productive in R. This is mostly true, in my experience.

1.1 The Mighty Data Frame

R comes with with a variety of built-in data sets that are very useful for getting used to data sets and how to manipulate them.

AirPassengers	Monthly Airline Passenger Numbers 1949–1960
BJsales	Sales Data with Leading Indicator
BOD	Biochemical Oxygen Demand



Figure 1.1:

C02	Carbon Dioxide Uptake in Grass Plants
ChickWeight	Weight versus age of chicks on different diets
DNase	Elisa assay of DNase
Formaldehyde	Determination of Formaldehyde
HairEyeColor	Hair and Eye Color of Statistics Students
Harman23.cor	Harman Example 2.3
Harman74.cor	Harman Example 7.4
Indometh	Pharmacokinetics of Indomethacin
InsectSprays	Effectiveness of Insect Sprays
JohnsonJohnson	Quarterly Earnings per Johnson & Johnson Share
LakeHuron	Level of Lake Huron 1875-1972
LifeCycleSavings	Intercountry Life-Cycle Savings Data
Loblolly	Growth of Loblolly pine trees
Nile	Flow of the River Nile
Orange	Growth of Orange Trees
OrchardSprays	Potency of Orchard Sprays
PlantGrowth	Results from an Experiment on Plant Growth
Puromycin	Reaction Velocity of an Enzymatic Reaction
Theoph	Pharmacokinetics of Theophylline

1.2 A Reference Data Frame

We will use a well-known data frame, at least in R circles, called **mtcars** which is part of any default installation of R. It is a simple data set relating to, well, automobiles. This data frame has the distinction of being the most (ab)used data frame in R education.

The data was extracted from the 1974 Motor Trend US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973–74 models).

A data frame with 32 observations on 11 (numeric) variables.

```
[, 1]    mpg Miles/(US) gallon
[, 2]    cyl Number of cylinders
[, 3]    disp    Displacement (cu.in.)
[, 4]    hp    Gross horsepower
[, 5]    drat    Rear axle ratio
[, 6]    wt    Weight (1000 lbs)
[, 7]    qsec    1/4 mile time
[, 8]    vs    Engine (0 = V-shaped, 1 = straight)
[, 9]    am    Transmission (0 = automatic, 1 = manual)
[,10]    gear    Number of forward gears
[,11]    carb    Number of carburetors
```

1.3 Relation to dplyr

What you will discover is that the **dplyr** package, which itself is part of the much larger **tidyverse** package set, extends upon the idea of the basic R data frame in a way that some feel is superior. It depends on your point of view though the **tidyverse** has a lot of what I call a philosophic consistency in it which makes it **very** useful once you get some concepts in mind.

While you could start exclusively with **dplyr** and the **tidyverse** the world is still full of older code. Plus, many of the advantages of **dplyr** only become quite apparent when compared to the “older way” of doing things. So my recommendation is to know how to deal with data frames in base R while also spending time to learn the **dplyr** way of doing things.

Chapter 2

Digging In

Data frames look like an Excel Spreadsheet. The rows are observations and the columns are variables or “features” that represent some measurement or character-based description of a given observation. When viewed from the row point of view, the data can be heterogenous. When viewed as a column, the data is homogenous.


```
data(mtcars)
mtcars
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
## Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
## Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1

## Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
## Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
## Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
## Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
## Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
## Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

We can do this with this data such as make plots or create models:

```
plot(mpg ~ wt, data=mtcars)
```



Let's create a regression model. It doesn't take long to realize that most functions in R will use a data frame as input. This means that you will spend a lot of time working with data frames to get them into shape for use with modeling and visualization tools. In fact you will spend most of your time **importing, transforming, and cleaning**.

```
(mylm <- lm(mpg ~ ., data = mtcars))
```

```
##
## Call:
## lm(formula = mpg ~ ., data = mtcars)
##
## Coefficients:
```

```
## (Intercept)          cyl          disp          hp          drat
##    12.30337        -0.11144        0.01334       -0.02148        0.78711
##           wt          qsec          vs          am          gear
##   -3.71530         0.82104        0.31776        2.52023        0.65541
##           carb
##   -0.19942
```

There are some useful functions that help you understand the structure of a data frame. One of the most important ones is called the `str()` function which is short hand for **structure**.

2.1 Structure

```
str(mtcars)

## 'data.frame':   32 obs. of  11 variables:
## $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
## $ cyl : num   6 6 4 6 8 6 8 4 4 6 ...
## $ disp: num  160 160 108 258 360 ...
## $ hp  : num  110 110 93 110 175 105 245 62 95 123 ...
## $ drat: num   3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
## $ wt  : num   2.62 2.88 2.32 3.21 3.44 ...
## $ qsec: num   16.5 17 18.6 19.4 17 ...
## $ vs  : num   0 0 1 1 0 1 0 1 1 1 ...
## $ am  : num   1 1 1 0 0 0 0 0 0 0 ...
## $ gear: num   4 4 4 3 3 3 3 4 4 4 ...
## $ carb: num   4 4 1 1 2 1 4 2 2 4 ...
```

This gives you some idea about the number of rows and columns of the data frame along with a description of the variable types and their values. I use this function frequently. Other functions that will help you include the following.

2.2 Meta Information

```
# how many rows
nrow(mtcars)

## [1] 32

# how many columns
ncol(mtcars)

## [1] 11
```

```
# Column names
names(mtcars)

## [1] "mpg" "cyl" "disp" "hp" "drat" "wt" "qsec" "vs" "am" "gear"
## [11] "carb"
```

2.3 Printing

Some data frames, such as `mtcars`, don't have many rows but others might have hundreds, thousands or even more than that ! Imagine trying to view one of those data frames. It is for this reason that the `head()` and `tail()` functions exist.

```
head(mtcars,5) # First 5 rows

##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110  3.90 2.620 16.46 0  1    4    4
## Mazda RX4 Wag  21.0   6  160 110  3.90 2.875 17.02 0  1    4    4
## Datsun 710      22.8   4  108  93  3.85 2.320 18.61 1  1    4    1
## Hornet 4 Drive  21.4   6  258 110  3.08 3.215 19.44 1  0    3    1
## Hornet Sportabout 18.7   8  360 175  3.15 3.440 17.02 0  0    3    2

tail(mtcars,3) # Last 3 rows

##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Ferrari Dino  19.7   6  145 175  3.62 2.77 15.5  0  1    5    6
## Maserati Bora 15.0   8  301 335  3.54 3.57 14.6  0  1    5    8
## Volvo 142E    21.4   4  121 109  4.11 2.78 18.6  1  1    4    2
```

2.4 Accessing Rows And Columns

There are various ways to select, remove, or exclude rows and columns of a data frame. We use the **bracket** notation to do this. This is very powerful. Keep in mind that data frames have rows and columns so it would make sense that you need a way to specify what rows and columns you want to access.

```
mtcars[1,] # First row, all columns

##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4  21    6  160 110  3.9 2.62 16.46 0  1    4    4

mtcars[1:3,] # First three rows, all columns

##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4  21.0   6  160 110  3.90 2.620 16.46 0  1    4    4
## Mazda RX4 Wag 21.0   6  160 110  3.90 2.875 17.02 0  1    4    4
```

```
## Datsun 710      22.8   4  108  93 3.85 2.320 18.61  1  1    4    1
```

```
# All rows, and first 4 columns
```

```
mtcars[,1:4]
```

```
##              mpg cyl  disp  hp
## Mazda RX4      21.0   6 160.0 110
## Mazda RX4 Wag  21.0   6 160.0 110
## Datsun 710     22.8   4 108.0  93
## Hornet 4 Drive  21.4   6 258.0 110
## Hornet Sportabout 18.7   8 360.0 175
## Valiant        18.1   6 225.0 105
## Duster 360     14.3   8 360.0 245
## Merc 240D      24.4   4 146.7  62
## Merc 230       22.8   4 140.8  95
## Merc 280       19.2   6 167.6 123
## Merc 280C      17.8   6 167.6 123
## Merc 450SE     16.4   8 275.8 180
## Merc 450SL     17.3   8 275.8 180
## Merc 450SLC    15.2   8 275.8 180
## Cadillac Fleetwood 10.4   8 472.0 205
## Lincoln Continental 10.4   8 460.0 215
## Chrysler Imperial 14.7   8 440.0 230
## Fiat 128       32.4   4  78.7  66
## Honda Civic    30.4   4  75.7  52
## Toyota Corolla 33.9   4  71.1  65
## Toyota Corona  21.5   4 120.1  97
## Dodge Challenger 15.5   8 318.0 150
## AMC Javelin    15.2   8 304.0 150
## Camaro Z28     13.3   8 350.0 245
## Pontiac Firebird 19.2   8 400.0 175
## Fiat X1-9      27.3   4  79.0  66
## Porsche 914-2  26.0   4 120.3  91
## Lotus Europa   30.4   4  95.1 113
## Ford Pantera L 15.8   8 351.0 264
## Ferrari Dino   19.7   6 145.0 175
## Maserati Bora   15.0   8 301.0 335
## Volvo 142E     21.4   4 121.0 109
```

```
# Rows 1-5 and columns 1,2 and 8-10
```

```
mtcars[1:4,c(1:2,8:10)]
```

```
##              mpg cyl vs am gear
## Mazda RX4      21.0   6  0  1    4
## Mazda RX4 Wag  21.0   6  0  1    4
## Datsun 710     22.8   4  1  1    4
## Hornet 4 Drive  21.4   6  1  0    3
```

```
# Rows 1-5 and columns 1,2 and 8-10
mtcars[1:4,c(1:2,8:10)]
```

```
##           mpg cyl vs am gear
## Mazda RX4      21.0   6  0  1    4
## Mazda RX4 Wag  21.0   6  0  1    4
## Datsun 710      22.8   4  1  1    4
## Hornet 4 Drive 21.4   6  1  0    3
```

```
# Rows 1-5 and columns by name
mtcars[1:4,c("mpg","wt","drat")]
```

```
##           mpg    wt drat
## Mazda RX4      21.0 2.620 3.90
## Mazda RX4 Wag  21.0 2.875 3.90
## Datsun 710      22.8 2.320 3.85
## Hornet 4 Drive 21.4 3.215 3.08
```

2.5 Interrogation

Many times you will wish to find rows that satisfy certain conditions. For example, what rows have an `mpg > 11` and at `wt < 2.0`? We use the bracket notation to help us. We can pass logical conditions into the brackets. Note the following:

```
mtcars$mpg > 11 & mtcars$wt < 2.0
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [12] FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE FALSE FALSE
## [23] FALSE FALSE FALSE TRUE FALSE TRUE FALSE FALSE FALSE FALSE
```

There are 32 elements in this logical vector each with a value of either TRUE or FALSE. When passed into the row index of the bracket notation, it will print that row if the corresponding value is TRUE. If FALSE, the row will not be printed.

```
mtcars[mtcars$mpg > 11 & mtcars$wt < 2.0,]
```

```
##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Honda Civic  30.4   4  75.7  52 4.93 1.615 18.52  1  1   4     2
## Toyota Corolla 33.9   4  71.1  65 4.22 1.835 19.90  1  1   4     1
## Fiat X1-9     27.3   4  79.0  66 4.08 1.935 18.90  1  1   4     1
## Lotus Europa  30.4   4  95.1 113 3.77 1.513 16.90  1  1   5     2
```

What if we just want to know how many cars satisfy this condition?

```
nrow(mtcars[mtcars$mpg > 11 & mtcars$wt < 2.0,])
```

```
## [1] 4
```

Find all rows that correspond to cars with 4 cylinders

```
mtcars[mtcars$cyl == 4,]
```

```
##           mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb
## Datsun 710   22.8   4 108.0  93 3.85 2.320 18.61 1   1    4    1
## Merc 240D   24.4   4 146.7  62 3.69 3.190 20.00 1   0    4    2
## Merc 230    22.8   4 140.8  95 3.92 3.150 22.90 1   0    4    2
## Fiat 128    32.4   4  78.7  66 4.08 2.200 19.47 1   1    4    1
## Honda Civic 30.4   4  75.7  52 4.93 1.615 18.52 1   1    4    2
## Toyota Corolla 33.9  4  71.1  65 4.22 1.835 19.90 1   1    4    1
## Toyota Corona 21.5  4 120.1  97 3.70 2.465 20.01 1   0    3    1
## Fiat X1-9    27.3   4  79.0  66 4.08 1.935 18.90 1   1    4    1
## Porsche 914-2 26.0  4 120.3  91 4.43 2.140 16.70 0   1    5    2
## Lotus Europa 30.4   4  95.1 113 3.77 1.513 16.90 1   1    5    2
## Volvo 142E  21.4   4 121.0 109 4.11 2.780 18.60 1   1    4    2
```

We can even use other R functions in the bracket notation. Extract all rows whose MPG value exceeds the mean MPG for the entire data frame.

```
mtcars[mtcars$mpg > mean(mtcars$mpg),]
```

```
##           mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb
## Mazda RX4   21.0   6 160.0 110 3.90 2.620 16.46 0   1    4    4
## Mazda RX4 Wag 21.0   6 160.0 110 3.90 2.875 17.02 0   1    4    4
## Datsun 710   22.8   4 108.0  93 3.85 2.320 18.61 1   1    4    1
## Hornet 4 Drive 21.4   6 258.0 110 3.08 3.215 19.44 1   0    3    1
## Merc 240D   24.4   4 146.7  62 3.69 3.190 20.00 1   0    4    2
## Merc 230    22.8   4 140.8  95 3.92 3.150 22.90 1   0    4    2
## Fiat 128    32.4   4  78.7  66 4.08 2.200 19.47 1   1    4    1
## Honda Civic 30.4   4  75.7  52 4.93 1.615 18.52 1   1    4    2
## Toyota Corolla 33.9  4  71.1  65 4.22 1.835 19.90 1   1    4    1
## Toyota Corona 21.5  4 120.1  97 3.70 2.465 20.01 1   0    3    1
## Fiat X1-9    27.3   4  79.0  66 4.08 1.935 18.90 1   1    4    1
## Porsche 914-2 26.0  4 120.3  91 4.43 2.140 16.70 0   1    5    2
## Lotus Europa 30.4   4  95.1 113 3.77 1.513 16.90 1   1    5    2
## Volvo 142E  21.4   4 121.0 109 4.11 2.780 18.60 1   1    4    2
```

Now find the cars for which the MPG exceeds the 75% percentile value for MPG

```
mtcars[mtcars$mpg > quantile(mtcars$mpg)[4],]
```

```
##           mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb
## Merc 240D   24.4   4 146.7  62 3.69 3.190 20.00 1   0    4    2
## Fiat 128    32.4   4  78.7  66 4.08 2.200 19.47 1   1    4    1
```

```
## Honda Civic      30.4   4  75.7  52 4.93 1.615 18.52  1  1   4   2
## Toyota Corolla  33.9   4  71.1  65 4.22 1.835 19.90  1  1   4   1
## Fiat X1-9        27.3   4  79.0  66 4.08 1.935 18.90  1  1   4   1
## Porsche 914-2    26.0   4 120.3  91 4.43 2.140 16.70  0  1   5   2
## Lotus Europa     30.4   4  95.1 113 3.77 1.513 16.90  1  1   5   2
```

2.6 Missing values

This is big deal. Most “real” data has rows that do not contain values for all columns. This is the so called “missing value” problem. Here is an example. The following code will read in a version of the mtcars data frame that has some missing values:

```
url <- "https://raw.githubusercontent.com/stevie42/utilities/master/data/mtcars_na.csv"
(mtcars_na <- read.csv(url, stringsAsFactors = FALSE))
```

```
##      mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## 1  21.0   6  160.0  110 3.90 2.620 16.46  0  1   4    4
## 2  21.0   6  160.0  110 3.90   NA 17.02  0  1   4    4
## 3  22.8   4  108.0   93 3.85 2.320 18.61  1  1   4    1
## 4  21.4   6  258.0  110 3.08 3.215 19.44  1  0   3    1
## 5  18.7   8  360.0  175 3.15 3.440 17.02  0  0   3    2
## 6  18.1   6  225.0  105 2.76 3.460 20.22  1  0   3    1
## 7  14.3   8  360.0  245 3.21 3.570 15.84  0  0   3    4
## 8  24.4   4  146.7   62 3.69 3.190 20.00  1  0   4    2
## 9  22.8   4  140.8   95 3.92   NA 22.90  1  0   4    2
## 10 19.2   6  167.6  123 3.92 3.440 18.30  1  0   4   NA
## 11 17.8   6  167.6  123 3.92 3.440 18.90  1  0   4    4
## 12 16.4   8  275.8  180 3.07 4.070 17.40  0  0   3   NA
## 13 17.3   8  275.8  180 3.07 3.730 17.60  0  0   3    3
## 14 15.2   8  275.8  180 3.07 3.780 18.00  0  0   3    3
## 15 10.4   8  472.0  205 2.93 5.250 17.98  0  0   3    4
## 16 10.4   8  460.0  215 3.00 5.424 17.82  0  0   3    4
## 17 14.7   8  440.0  230 3.23 5.345 17.42  0  0   3    4
## 18 32.4   4   78.7   66 4.08 2.200 19.47  1  1   4    1
## 19 30.4   4   75.7   52 4.93 1.615 18.52  1  1   4   NA
## 20 33.9   4   71.1   65 4.22 1.835 19.90  1  1   4   NA
## 21 21.5   4  120.1   97 3.70 2.465 20.01  1  0   3    1
## 22 15.5   8  318.0  150 2.76 3.520 16.87  0  0   3    2
## 23 15.2   8  304.0  150 3.15   NA 17.30  0  0   3   NA
## 24 13.3   8  350.0  245 3.73 3.840 15.41  0  0   3    4
## 25 19.2   8  400.0  175 3.08 3.845 17.05  0  0   3    2
## 26 27.3   4   79.0   66 4.08 1.935 18.90  1  1   4    1
## 27 26.0   4  120.3   91 4.43 2.140 16.70  0  1   5    2
## 28 30.4   4   95.1  113 3.77 1.513 16.90  1  1   5   NA
```



```
## 29 15.8   8 351.0 264 4.22 3.170 14.50 0 1   5   4
## 30 19.7   6 145.0 175 3.62 2.770 15.50 0 1   5   6
## 31 15.0   8 301.0 335 3.54 3.570 14.60 0 1   5   8
## 32 21.4   4 121.0 109 4.11 2.780 18.60 1 1   4   2
```

If you look, you can see the missing values “NA” present in certain columns. This is R’s way of indicating what is missing. There are functions that can help you find these. This is important because, for example, if you wanted to find the average value of a column, say the `wt` column then there will be a problem as it contains a missing value:

```
mean(mtcars_na$wt)
```

```
## [1] NA
```

We have to tell the function to remove missing values from consideration.

```
mean(mtcars$wt, na.rm=TRUE)
```

```
## [1] 3.21725
```

A more general approach would involve the following functions.

```
complete.cases(mtcars_na)
```

```
## [1] TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE TRUE
## [12] FALSE TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE TRUE TRUE
## [23] FALSE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE
```

```
# How many rows in the df do not contain any NAs ?
```

```
sum(complete.cases(mtcars_na))
```

```
## [1] 24
```

```
# How many rows in the df do contain at least one NA ?
```

```
sum(!complete.cases(mtcars_na))
```

```
## [1] 8
```

How would we find those rows and print them ?

```
mtcars_na[complete.cases(mtcars_na),]
```

```
##      mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## 1  21.0   6  160.0 110 3.90 2.620 16.46 0 1   4   4
## 3  22.8   4  108.0  93 3.85 2.320 18.61 1 1   4   1
## 4  21.4   6  258.0 110 3.08 3.215 19.44 1 0   3   1
## 5  18.7   8  360.0 175 3.15 3.440 17.02 0 0   3   2
## 6  18.1   6  225.0 105 2.76 3.460 20.22 1 0   3   1
## 7  14.3   8  360.0 245 3.21 3.570 15.84 0 0   3   4
## 8  24.4   4  146.7  62 3.69 3.190 20.00 1 0   4   2
## 11 17.8   6  167.6 123 3.92 3.440 18.90 1 0   4   4
```

```
## 13 17.3 8 275.8 180 3.07 3.730 17.60 0 0 3 3
## 14 15.2 8 275.8 180 3.07 3.780 18.00 0 0 3 3
## 15 10.4 8 472.0 205 2.93 5.250 17.98 0 0 3 4
## 16 10.4 8 460.0 215 3.00 5.424 17.82 0 0 3 4
## 17 14.7 8 440.0 230 3.23 5.345 17.42 0 0 3 4
## 18 32.4 4 78.7 66 4.08 2.200 19.47 1 1 4 1
## 21 21.5 4 120.1 97 3.70 2.465 20.01 1 0 3 1
## 22 15.5 8 318.0 150 2.76 3.520 16.87 0 0 3 2
## 24 13.3 8 350.0 245 3.73 3.840 15.41 0 0 3 4
## 25 19.2 8 400.0 175 3.08 3.845 17.05 0 0 3 2
## 26 27.3 4 79.0 66 4.08 1.935 18.90 1 1 4 1
## 27 26.0 4 120.3 91 4.43 2.140 16.70 0 1 5 2
## 29 15.8 8 351.0 264 4.22 3.170 14.50 0 1 5 4
## 30 19.7 6 145.0 175 3.62 2.770 15.50 0 1 5 6
## 31 15.0 8 301.0 335 3.54 3.570 14.60 0 1 5 8
## 32 21.4 4 121.0 109 4.11 2.780 18.60 1 1 4 2
```

And here are the ones that do contain missing values:

```
mtcars_na[!complete.cases(mtcars_na),]
```

```
##      mpg cyl  disp  hp drat    wt  qsec vs am gear carb
## 2  21.0   6 160.0 110 3.90    NA 17.02  0  1   4    4
## 9  22.8   4 140.8  95 3.92    NA 22.90  1  0   4    2
## 10 19.2   6 167.6 123 3.92  3.440 18.30  1  0   4    NA
## 12 16.4   8 275.8 180 3.07  4.070 17.40  0  0   3    NA
## 19 30.4   4  75.7  52 4.93  1.615 18.52  1  1   4    NA
## 20 33.9   4  71.1  65 4.22  1.835 19.90  1  1   4    NA
## 23 15.2   8 304.0 150 3.15    NA 17.30  0  0   3    NA
## 28 30.4   4  95.1 113 3.77  1.513 16.90  1  1   5    NA
```

One quick way to omit rows with missing values is:

```
na.omit(mtcars_na)
```

```
##      mpg cyl  disp  hp drat    wt  qsec vs am gear carb
## 1  21.0   6 160.0 110 3.90  2.620 16.46  0  1   4    4
## 3  22.8   4 108.0  93 3.85  2.320 18.61  1  1   4    1
## 4  21.4   6 258.0 110 3.08  3.215 19.44  1  0   3    1
## 5  18.7   8 360.0 175 3.15  3.440 17.02  0  0   3    2
## 6  18.1   6 225.0 105 2.76  3.460 20.22  1  0   3    1
## 7  14.3   8 360.0 245 3.21  3.570 15.84  0  0   3    4
## 8  24.4   4 146.7  62 3.69  3.190 20.00  1  0   4    2
## 11 17.8   6 167.6 123 3.92  3.440 18.90  1  0   4    4
## 13 17.3   8 275.8 180 3.07  3.730 17.60  0  0   3    3
## 14 15.2   8 275.8 180 3.07  3.780 18.00  0  0   3    3
## 15 10.4   8 472.0 205 2.93  5.250 17.98  0  0   3    4
## 16 10.4   8 460.0 215 3.00  5.424 17.82  0  0   3    4
```

```
## 17 14.7    8 440.0 230 3.23 5.345 17.42 0 0    3    4
## 18 32.4    4  78.7  66 4.08 2.200 19.47 1 1    4    1
## 21 21.5    4 120.1  97 3.70 2.465 20.01 1 0    3    1
## 22 15.5    8 318.0 150 2.76 3.520 16.87 0 0    3    2
## 24 13.3    8 350.0 245 3.73 3.840 15.41 0 0    3    4
## 25 19.2    8 400.0 175 3.08 3.845 17.05 0 0    3    2
## 26 27.3    4  79.0  66 4.08 1.935 18.90 1 1    4    1
## 27 26.0    4 120.3  91 4.43 2.140 16.70 0 1    5    2
## 29 15.8    8 351.0 264 4.22 3.170 14.50 0 1    5    4
## 30 19.7    6 145.0 175 3.62 2.770 15.50 0 1    5    6
## 31 15.0    8 301.0 335 3.54 3.570 14.60 0 1    5    8
## 32 21.4    4 121.0 109 4.11 2.780 18.60 1 1    4    2
```

2.7 Continuous vs Factors

One **recipe** that I use frequently is given below. This tells me how many unique values are assumed by each column which then helps to identify continuous quantities and categories. If a column assumes only a small number of unique values then perhaps it should be classified as a factor. Don't let the code here scare you. If you are new to R and don't yet understand what is going on then just use this as a "recipe" for now.

```
sapply(mtcars, function(x) length(unique(x)))
```

```
## mpg  cyl disp  hp drat   wt  qsec    vs  am gear carb
##   25    3  27   22  22   29   30     2   2   3    6
```

So it looks to me, for example, that **cyl**, **vs**, **am**, **gear**, and **carb** are actually categories rather than measured quantities. If you look at the help page for **mtcars** you will see that **am** is a 0 or 1 which corresponds to, respectively, a car with an automatic transmission (0) or a manual transmission (1). If you use the **summarize** function it will treat this variable as a numeric, continuous quantity.

Is it actually possible to have a transmission value of 0.4062 ?

```
summary(mtcars$am)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000  0.0000  0.0000  0.4062  1.0000  1.0000
```

I might then use some code to transform this into factors so that when they are used with various modeling functions they will be recognized as such. For example, if we summarize the data frame right now, we will see the following

```
summary(mtcars$am)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
## 0.0000 0.0000 0.0000 0.4062 1.0000 1.0000
```

Let's turn **am** into a factor

```
mtcars$am <- factor(mtcars$am,
                    levels = c(0,1),
                    labels = c("Auto","Man") )
```


Now the summary will make more sense. This is also useful because graphics packages such as ggplot2 will know how to handle factors.

```
summary(mtcars$am)
```

```
## Auto  Man
```

```
##   19   13
```

```
ggplot(mtcars,aes(x=wt,y=mpg)) +
  geom_point() +
  facet_wrap(~am)
```



2.8 Sorting data

Sorting rows in a data frame is a common activity. However, in Base R this is called “ordering” because of the function used to “order” the data. Let's say we want to sort or “order” the mtcars data frame such that the row with the

lowest mpg value is listed first and the row with the highest mpg value is listed last. First, look at the **order** function's output. What are those numbers ?

```
order(mtcars$mpg)
```

```
## [1] 15 16 24 7 17 31 14 23 22 29 12 13 11 6 5 10 25 30 1 2 4 32 21
## [24] 3 9 8 27 26 19 28 18 20
```

Oh, so they are row numbers corresponding to rows in `mtcars`. Row 15 has the car with the lowest mpg. Row 16 corresponds to the car with the next lowest mpg and so on. So we can use this information to order our dataframe accordingly:

```
mtcars[order(mtcars$mpg),]
```

```
##          mpg cyl  disp  hp drat   wt  qsec vs  am gear carb
## Cadillac Fleetwood 10.4  8 472.0 205 2.93 5.250 17.98 0 Auto    3    4
## Lincoln Continental 10.4  8 460.0 215 3.00 5.424 17.82 0 Auto    3    4
## Camaro Z28         13.3  8 350.0 245 3.73 3.840 15.41 0 Auto    3    4
## Duster 360         14.3  8 360.0 245 3.21 3.570 15.84 0 Auto    3    4
## Chrysler Imperial  14.7  8 440.0 230 3.23 5.345 17.42 0 Auto    3    4
## Maserati Bora       15.0  8 301.0 335 3.54 3.570 14.60 0 Man     5    8
## Merc 450SLC        15.2  8 275.8 180 3.07 3.780 18.00 0 Auto    3    3
## AMC Javelin        15.2  8 304.0 150 3.15 3.435 17.30 0 Auto    3    2
## Dodge Challenger   15.5  8 318.0 150 2.76 3.520 16.87 0 Auto    3    2
## Ford Pantera L     15.8  8 351.0 264 4.22 3.170 14.50 0 Man     5    4
## Merc 450SE         16.4  8 275.8 180 3.07 4.070 17.40 0 Auto    3    3
## Merc 450SL         17.3  8 275.8 180 3.07 3.730 17.60 0 Auto    3    3
## Merc 280C          17.8  6 167.6 123 3.92 3.440 18.90 1 Auto    4    4
## Valiant            18.1  6 225.0 105 2.76 3.460 20.22 1 Auto    3    1
## Hornet Sportabout  18.7  8 360.0 175 3.15 3.440 17.02 0 Auto    3    2
## Merc 280           19.2  6 167.6 123 3.92 3.440 18.30 1 Auto    4    4
## Pontiac Firebird   19.2  8 400.0 175 3.08 3.845 17.05 0 Auto    3    2
## Ferrari Dino       19.7  6 145.0 175 3.62 2.770 15.50 0 Man     5    6
## Mazda RX4          21.0  6 160.0 110 3.90 2.620 16.46 0 Man     4    4
## Mazda RX4 Wag      21.0  6 160.0 110 3.90 2.875 17.02 0 Man     4    4
## Hornet 4 Drive     21.4  6 258.0 110 3.08 3.215 19.44 1 Auto    3    1
## Volvo 142E         21.4  4 121.0 109 4.11 2.780 18.60 1 Man     4    2
## Toyota Corona      21.5  4 120.1  97 3.70 2.465 20.01 1 Auto    3    1
## Datsun 710         22.8  4 108.0  93 3.85 2.320 18.61 1 Man     4    1
## Merc 230           22.8  4 140.8  95 3.92 3.150 22.90 1 Auto    4    2
## Merc 240D          24.4  4 146.7  62 3.69 3.190 20.00 1 Auto    4    2
## Porsche 914-2      26.0  4 120.3  91 4.43 2.140 16.70 0 Man     5    2
## Fiat X1-9          27.3  4  79.0  66 4.08 1.935 18.90 1 Man     4    1
## Honda Civic        30.4  4  75.7  52 4.93 1.615 18.52 1 Man     4    2
## Lotus Europa       30.4  4  95.1 113 3.77 1.513 16.90 1 Man     5    2
## Fiat 128           32.4  4  78.7  66 4.08 2.200 19.47 1 Man     4    1
```

```
## Toyota Corolla      33.9   4  71.1  65 4.22 1.835 19.90  1  Man    4    1
```

To invert the sense of the order use the **rev** function. We'll also use the **head** function to list only the first 5 rows of the result. Note that in base R, using composite functions is welcomed although you will find out that this is not a value in the tidyverse. For math people, using a composite function is natural which, in large part, is why R embraced that approach early on.

```
head(mtcars[rev(order(mtcars$mpg)),])
```

```
##           mpg cyl  disp  hp drat   wt  qsec vs  am  gear carb
## Toyota Corolla 33.9   4  71.1  65 4.22 1.835 19.90  1  Man    4    1
## Fiat 128       32.4   4  78.7  66 4.08 2.200 19.47  1  Man    4    1
## Lotus Europa   30.4   4  95.1 113 3.77 1.513 16.90  1  Man    5    2
## Honda Civic    30.4   4  75.7  52 4.93 1.615 18.52  1  Man    4    2
## Fiat X1-9      27.3   4  79.0  66 4.08 1.935 18.90  1  Man    4    1
## Porsche 914-2  26.0   4 120.3  91 4.43 2.140 16.70  0  Man    5    2
```

2.9 Creating Data Frames From CSV Files

Many times data will be read in from a comma delimited file exported from Excel. These are known as Comma Separated Value files - generally abbreviated as **CSV**. The file can be read from a local drive or even from the Web as long as you know the URL associated with the file. In this example, there is a file on the Internet relating to some testing data involving students and various subjects.

```
url <- "https://raw.githubusercontent.com/pittardsp/bios545r_spring_2018/master/SUPPORTING_FILES/01%20Data%20Frames/01%20Data%20Frames.csv"
```

```
data1 <- read.csv(url,header=T,sep=",")
```

```
head(data1)
```

```
##    id female race ses schtyp prog read write math science socst
## 1  70      0    4   1     1    1   57    52   41      47     57
## 2 121      1    4   2     1    3   68    59   53      63     61
## 3  86      0    4   3     1    1   44    33   54      58     31
## 4 141      0    4   3     1    3   63    44   47      53     56
## 5 172      0    4   2     1    2   47    52   57      53     61
## 6 113      0    4   2     1    2   44    52   51      63     61
```

Chapter 3

Literature

Here is a review of existing methods.

Chapter 4

Methods

We describe our methods in this chapter.